Non-linear least Squares

unknown true model: $y^* = f_{\beta^*}(x) + \varepsilon$ $\quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$

$f_{\beta^*}(x)$: non-linear fct with true parameters β^*

[special case: $f_{\beta^*}(x) = x \cdot \beta^*$ (linear fct.) $\Rightarrow$ ordinary least squares]

goal: $\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|y - f_\beta(x)\|^2$, in general no analytic solution

many approximations:

- regression trees, regression forests
- Levenberg-Marquardt algorithm
- change of variables to transform non-linear prob. into a linear one (approximately)
- kernel methods
- Gaussian processes
- neural networks
- ...

EML

Advanced Machine Learning

Regression Trees

work like density and decision trees: partition the feature space into bins using a tree, have a constant response for every bin:

$$\hat{y}_i = \hat{f}(x_i) = \sum_{\ell=1}^{L} f_\ell \, \mathbb{1}[x_i \in \text{bin}_\ell]$$

$\mathbb{1}[\ldots]$ ← true to find the bin for x_i

$\ell = 1$ ↑ sum over all bins; f_ℓ ↑ constant bin response

find optimal response in bin_ℓ via least squares loss (assume bins are already fixed):

$$\hat{f}_\ell = \int_{x \in \text{bin}_\ell} (f_\ell - f(x))^2 \, dx \qquad \hat{f}(x) = \sum_{\ell=1}^{L} f_\ell \, \mathbb{1}[x_i \in \text{bin}_\ell]$$

$$\{\hat{f}_\ell\} = \arg\min_{\{f_\ell\}} \int (\hat{f}(x) - f(x))^2 \, p(x) \, dx$$

$$= \arg\min_{\{f_\ell\}} \underbrace{\int f(x)^2 p(x) \, dx}_{\text{independent of model} \Rightarrow \text{drop}} - 2 \int \hat{f}(x) f(x) p(x) \, dx + \int \hat{f}(x)^2 p(x) \, dx$$

$$= \arg\min_{\{f_\ell\}} -2 \sum_{\ell=1} f_\ell \int_{x \in \text{bin}_\ell} f(x) p(x) \, dx + \sum_\ell f_\ell^2 \int_{x \in \text{bin}_\ell} p(x) \, dx$$

$$\frac{\partial \text{loss}}{\partial f_\ell} = -2 \underbrace{\int_{x \in \text{bin}_\ell} f(x) p(x) \, dx}_{\mathbb{E}_{\text{bin}_\ell}[f(x)]} + 2 f_\ell \underbrace{\int_{x \in \text{bin}_\ell} p(x) \, dx}_{\text{total density in bin}_\ell} \doteq 0$$

$$\approx \left(\frac{1}{N_\ell} \sum_{i : x_i \in \text{bin}_\ell} f(x_i) \right) \cdot \frac{N_\ell}{N} + f_\ell \frac{N_\ell}{N} \doteq 0$$

$$\hat{f}_\ell = \frac{1}{N_\ell} \sum_{i : x_i \in \text{bin}_\ell} \underbrace{f(x_i)}_{= y_i}$$

$$\boxed{\hat{f}_\ell = \frac{1}{N_\ell} \sum_{i \in \text{bin}_\ell} y_i}$$

regression tree: piece-wise constant approximation of a non-linear function

truth: $Y = f(X) + \epsilon$ approx: $Y = \hat{f}(x) = \sum_{\ell=1}^{L} \hat{f}_\ell \, \mathbb{1}[X \in \text{bin}_\ell]$

minimize the squared loss

$$\{\hat{f}_\ell\}_{\ell=1}^{L} = \underset{\{f_\ell\}}{\operatorname{argmin}} \; \mathbb{E}_{p(x)}\left[(f(x) - \hat{f}(x))^2\right] = \underset{\{f_\ell\}}{\operatorname{argmin}} \int (f(x) - \hat{f}(x))^2 p(x)\, dx$$

$$= \underset{\{f_\ell\}}{\operatorname{argmin}} \; -2 \sum_{\ell=1}^{L} f_\ell \underbrace{\int_{x \in \text{bin}_\ell} f(x)\, p(x)\, dx}_{\sum_{i: x_i \in \text{bin}_\ell} y_i \; (= f(x_i)) \; \frac{1}{N}} + \sum_{\ell=1}^{L} f_\ell^2 \underbrace{\int_{x \in \text{bin}_\ell} 1 \cdot p(x)\, dx}_{\sum_{i: x_i \in \text{bin}_\ell} 1 \cdot \frac{1}{N}}$$

$$= \underset{\{f_\ell\}}{\operatorname{argmin}} \; -2 \sum_{\ell=1}^{L} f_\ell \cdot \frac{1}{N} \sum_{i: x_i \in \text{bin}_\ell} y_i + \sum_{\ell=1}^{L} f_\ell^2 \frac{N_\ell}{N}$$

$$\frac{\partial \text{Loss}}{\partial f_\ell} = -2 \frac{1}{N} \sum_{i: x_i \in \text{bin}_\ell} y_i + 2 f_\ell \frac{N_\ell}{N} \stackrel{!}{=} 0$$

$$\boxed{\hat{f}_\ell = \frac{1}{N_\ell} \sum_{i: x_i \in \text{bin}_\ell} y_i}$$

training alg.: exactly like decision & dunkity tree, just with different criterion for selecting the optimal split: in each node to be split, iterate over all candidate splits and compute the squared error: $f_{\text{left}} = \frac{1}{N_{\text{left}}} \sum_{i: x_i \in \text{left}} y_i$ $f_{\text{right}} = \frac{1}{N_{\text{right}}} \sum_{i: x_i \in \text{right}} y_i$

execute split with minimal R: $R = \sum_{i: x_i \in \text{left}} (y_i - f_{\text{left}})^2 + \sum_{i: x_i \in \text{right}} (y_i - f_{\text{right}})^2$

regression forest: to avoid overfitting, train an ensemble of regression trees with the usual tricks (each tree is trained with a bootstrap sample of the TS, in each node a random subset of the features is considered for splitting) and return the average response of all trees

Levenberg - Marquardt - Algorithm

setting here: $Y = f_\beta(X) + \varepsilon$ with $f_\beta(X)$ a non-linear function with parameters β, functional form (formula) of f_β is known, only the true parameters β^* are unknown

[notice: linear models $f_\beta(X) = X \cdot \beta$ are a special case]

difficulty: - for $f_\beta(X)$ non-linear, there is usually no analytic solution for β

⇒ reduce to a sequence of linear problems via Taylor series, iteratively solve non-linear optimization by solving a least-squares prob in every iteration

least-squares objective: $\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^{N} (y_i - f_\beta(x_i))^2$

- iterative Gauss-Newton alg. want to minimize $L(\beta)$
  - given current guess $\beta^{(t-1)}$, expand $L(\beta)$ into Taylor series at $\beta^{(t-1)}$

$$L(\underbrace{\beta^{(t-1)} + \Delta}_{=\beta^{(t)}}) \approx L(\beta^{(t-1)}) + \underbrace{\frac{\partial L}{\partial \beta}\Big|_{\beta^{(t-1)}}}_{\text{"Jacobian of the loss"}} \cdot \Delta + \text{higher orders (ignored)}$$

- least-squares minimization w.r.t. improvement Δ

$$\hat{\Delta} = \arg\min_{\Delta} \left(\underbrace{L(\beta^{(t-1)})}_{\tilde{Y}} + \underbrace{\frac{\partial L}{\partial \beta}\Big|_{\beta^{(t-1)}}}_{-\tilde{X}} \cdot \Delta \right)^2$$

auxilliary variables: $\tilde{Y}$ (vector of length N), $-\tilde{X}$ (matrix of derivatives $N \times D$)

$$\hat{\Delta} = \arg\min_{\Delta} \left(\tilde{Y} - \tilde{X} \cdot \Delta \right)^2 \qquad \hat{\Delta} = (\tilde{X}^T \tilde{X})^{-1} \cdot \tilde{X}^T \cdot \tilde{Y}$$

solution via pseudo-inverse

not yet optimal, because it tends to overfit $\Rightarrow$ L_2 regularization

$$\hat{\Delta} = \arg\min_{\Delta} \left(\tilde{Y} - \tilde{X} \cdot \Delta \right)^2 + \tau \|\Delta\|^2$$

- in case of non-linear least squares:

$$L(\beta) = Y - f_\beta(X)$$

$$L(\beta^{(t-1)} + \Delta) = \underbrace{L(\beta^{(t-1)})}_{= Y - f_{\beta^{(t-1)}}(X)} + \underbrace{\frac{\partial f_\beta(X)}{\partial \beta}\Big|_{\beta^{(t-1)}}}_{-\tilde{X}} \cdot \Delta$$

$$= \tilde{Y} - \tilde{X} \cdot \Delta$$

optimize via ridge-regression: $\hat{\Delta} = (\tilde{X}^T \tilde{X} + \tau \mathbb{I})^{-1} \tilde{X}^T \tilde{Y}$

$$\beta^{(t)} = \beta^{(t-1)} + \Delta$$

- why is it good to use ridge regression?

case 1: $\tilde{X}^T\tilde{X}$ has good condition $\Rightarrow \tau \approx 0$ (no regularization needed)

$\Rightarrow \hat{\Delta}$ is unbiased least-squares solution $\Rightarrow$ very fast convergence (few iterations)

case 2: $\tilde{X}^T\tilde{X}$ has bad condition $\Rightarrow \tau \gg 0$,

in the extreme, $\lambda^T\tilde{X}$ is dominated by τI

$$(\tilde{X}^T\tilde{X} + \tau I)^{-1}\tilde{X}^T\tilde{Y} \Rightarrow \frac{1}{\tau}\tilde{X}^T\tilde{Y}$$

$\Rightarrow$ we get a gradient descent alg. with learning rate $\frac{1}{2\tau}$

$$\frac{\partial(\tilde{X}-\tilde{X}\cdot\Delta)^2}{\partial\Delta}\Big|_{\Delta=0} = -2\tilde{X}^T\tilde{Y}$$

$\Rightarrow \tau$ allows us to interpolate between fast Newton iterations and slower, but numerically stable, gradient descent

- Questions: - how to choose good τ? $\Rightarrow$ self-regulating algorithm
  - how to ensure that τ has the same effect on all features

[recall: in ridge regression, we standardized the features $\tilde{x}_j = \frac{x_j - \bar{x}_j}{\mathrm{std}(x_j)}$]

standardization is now impossible, because $\tilde{X}$ is the Jacobian matrix $\left(\frac{\partial f}{\partial p}\right)$, which would become incorrect by scaling

$\Rightarrow$ Marquardt trick: weight τ with $\|\tilde{x}_j\|^2 = \mathrm{diag}(\tilde{X}^T\tilde{X})$

$$\hat{\Delta} = \left(\tilde{X}^T\tilde{X} + \tau\,\mathrm{diag}(\tilde{X}^T\tilde{X})\right)^{-1}\cdot\lambda^T\tilde{Y}$$

(in practice: compute $S = \tilde{X}^T\tilde{X}$, then scale diagonal elements $S_{jj} \Leftarrow S_{jj}(1+\tau)$)

full Levenberg - Marquardt algorithm

⓪ define initial guess $\beta^{(0)}$ (this is usually critical, bad guess ⇒ bad optimum)

initial regularization strength $\tau^{(0)}$

① for $t = 1, \dots, T$ (or until convergence):

ⓐ compute $\vec{y} = Y - f_{\beta^{(t-1)}}(X)$, $\tilde{X} = \left.\frac{\partial f_\beta(x)}{\partial \beta}\right|_{\beta^{(t-1)}}$

ⓑ solve $(\tilde{X}^T \tilde{X} + \tau^{(t-1)} \cdot \mathrm{diag}(\tilde{X}^T \tilde{X})) \cdot \Delta = \tilde{X}^T \cdot \vec{y}$ by some linear solver ⇒ $\vec{\Delta}$

ⓒ compute $\vec{\beta} = \beta^{(t-1)} + \vec{\Delta}$ and residuals $(Y - f_{\vec{\beta}}(X))^2$

ⓓ if residual (squared loss) of $\vec{\beta}$ is smaller than of $\beta^{(t-1)}$

- accept $\vec{\beta}$ as our new guess $\beta^{(t)} = \vec{\beta}$
- reduce regularization strength $\tau^{(t)} = \tau^{(t-1)} / \sqrt{2}$ (hyperparam)

otherwise

- reject the guess $\vec{\beta}$: $\beta^{(t)} = \beta^{(t-1)}$
- increase the regularization $\tau^{(t)} = \sqrt{2}\, \tau^{(t-1)}$

$$\mathrm{diag}(A) = \begin{bmatrix} A_{11} & & 0 \\ & A_{22} & \\ & & \ddots \\ 0 & & A_{NO} \end{bmatrix}$$